Research Computing New User Seminar

Be Boulder.

University of Colorado **Boulder**

Research Computing New User Seminar

- Joel Frahm
- Joel.Frahm@Colorado.Edu
- www.rc.colorado.edu
- Slides: https://github.com/ResearchComputing/New_User_Seminar

Before We Begin

- Goals
  - Inform people about RC resources, expectations, etc.
  - Reduce user frustrations, questions
  - Avoid misunderstandings, lost time, lost work
  - Inform users about best practices

- Things to take particular note of
  - Confusing, ambiguous, highly nuanced concepts
  - Common mistakes or frustrations
  - Best Practices

- Good questions to ask
  - Why? Questions
  - If a question is said to be covered later feel free to re-ask if it's not answered to your satisfaction.

Outline

- What is RC?
    - Resources and services/support
    - Summit overview
- Steps to get access to our systems
    - Accounts
    - Two-factor authentication
    - Logging in
- Navigating our systems
    - Blanca
    - Petalibrary
- Summit (and other RC cluster) New user
    - Allocations
    - Storage spaces
    - Data transfer - Globus
    - Software
- Running jobs

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

What is Research Computing?

- Provide services for researchers that include:
  - Large scale computing
  - Data storage
  - High speed data transfer
  - Data management support
  - Consulting
  - Training

- We are likely best known for:
  - Summit Supercomputer
  - PetaLibrary storage

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

What Would I Use Summit For?

- Research Computing is more than just Summit
- But it is what we are most known for
- So what would you use Summit For?
  - Solving large problems that require more:
    - Memory than you have on your personal computer
    - Cores/nodes/power than you have on your personal computer
  - High performance GPU computing
  - High memory jobs
  - Visualization rendering
- Not a place for:
  - Large data storage

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Hardware - Summit Supercomputer

- 450+ compute nodes (mostly Intel Xeon Haswell)

- 24 cores per "shas" (general compute) node, different core counts for other node types

- 11,400 total cores

- Omni-Path network

- 1.2 PB scratch storage

- GPFS Scratch File system

- 67% CU, 23% CSU, 10% RMACC

Additional Types of Summit Compute Nodes

- 10 Graphics Processing Unit (GPU) Nodes
  - NVIDIA Tesla K80 (2/node)
- 5 High Memory Nodes
  - 2 TB of memory/node, 48 cores/node
- Phi Nodes
  - 20 nodes
  - Intel Xeon Phi

How To Access RC Resources?

1. Get an account
2. Set up two factor authentication
3. Inform us of any specific needs
4. Log in
5. Create greatness

• After you login, you will need to do many additional things that we will discuss today

Getting an RC Account

- CU Boulder users and affiliates:
- Request an account through the RC Account request portal
    - https://portals.rc.colorado.edu/accounts/account-request/create
    - If you are requesting access to Summit make sure you indicate this

- CSU Users:
    - Request an CSU eID if you don't have one
    - Fill out account application form
    - Duo authentication
    - Then get an RC user account
    - https://www.acns.colostate.edu/hpc/summit-get-started/

- RMACC Users:
    - Contact your local representative, if known. Email rc-help@colorado.edu also
    - We'll guide you through the process

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Setting up Two-Factor Authentication

- Required to provide an extra level of authentication, we are outside the firewall!

- Duo
  - Most users use the Duo smartphone app
  - "Phone Call" or SMS (text) are alternatives
  - Physical code generator "token" available for $20

Duo Authentication

- Once you get an account, contact rc-help@colorado.edu to request a Duo invitation

- Once you get the invitation, you'll get a series of steps to complete Duo enrollment

- For UCB RC supports Duo "push", "SMS","phone call" and the physical "Token" (code generator) for authentication

- Greatly prefer "push" – it's the cheapest for UCB

- Once your Duo is set up UCB users can manage their Duo setup at duo.colorado.edu

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Logging In

- It's important to note that you are NOT logging into any specific resource, Summit, etc.

- When you log in, you land on our login nodes

- From there, you can access our other resources

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

RC Resource Access

- To login to an RC login node:

  `ssh username@login.rc.colorado.edu`

- If logging in with Duo, your enter your password as:

  `duo:identikey_password`

- If logging in with Vasco, your enter your password as:

  `Pin+six-digit number on Vasco`

Navigating our Systems

- Now that you've logged in, now what?
  - What are the different node types we have?
  - What are the different storage spaces?
    - What should I be putting in these storage spaces?
  - How do I transfer data around?
  - How do I deal with software?

Blanca

- If you are a Blanca user, you need an RC account, and we need to know what resources to connect you to.
- To run jobs as a Blanca user, once you've logged into a login node, load the Blanca slurm module

```
module load slurm/blanca
```

Use `--qos=blanca-<group-identifier>` for high priority access

`--qos=blanca` for low-priority access

- Only certain users have access to Blanca – paid service
- If you are unsure, you can ask your advisor or RC
  - But likely if you are unsure you don't have access

PetaLibrary

- To access the PetaLibrary, you login in to one our RC's login nodes as normal

- Then you cd to either /work/<groupname> or /archive/<groupname>, depending on your PetaLibrary service
  - <groupname> is the name set for your group when you set up the PetaLibrary service
  - You do not include the <>
- Only certain users have access to PetaLibrary – paid service
- If you are unsure, you can ask your advisor or RC
  - But likely if you are unsure you don't have access

Using Summit (and other RC clusters)

Allocations

- You will need a compute time allocation to use any of our compute resources.

- Blanca and Summit Condo allocations are part of the buy-in

- New users are offered a checkbox for Summit access when creating their account, this gets you added to the Summit General allocation.

- Otherwise, to request access please email rc-help@colorado.edu and ask for access to the General allocation
  - Need to provide a few sentences on your project

- Once you have some benchmarks, most users will want to get an allocation to get access to a lot larger "share" of Summit.

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Why Do I Need to request access to Summit?

- Not all UCB users need or want to use Summit

- I have an account – why do I need an allocation?
  - An account validates you are eligible to use RC compute resources
  - An allocation allows us to keep track of your use of the system
  - This is important because:
    - We need to make sure we have enough resources to accommodate all of our users
    - Helps for reporting to NSF and the CU Research & Innovation Office
  - Applying for an allocation beyond General:
    - Gives you higher priority in the system and access to more compute time

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

What is Fair Share?

- Fair share scheduling uses a complex formula to determine priority in queue

- Looks at load for each user and each QOS and balances utilization to fairly share resources
  - Involves historical use by user plus how long job has been in the queue

- System will first look at weighted average utilization of user mostly over the last 4 weeks

- Then compare it to the fair share target percentage of a user

Fair Share Target Percentage

- The target percentage depends on your priority based on your project proposal
- Everyone not associated with a project shares a target percentage of 13% (20% of the CU fraction)
  - No guaranteed level per user
- If you are under (over) your target percentage (based on a 4 week average) your priority is increased (decreased)
- Reminder: this all only impacts pending jobs
- If no other pending jobs and enough resources are available then your job will run regardless of your previous usage

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Allocations

- Did not request Summit access when you signed up?

- Make a request now!
  - Include 2-3 sentences describing your proposed usage
  - Email rc-help@colorado.edu

- If you need access to a condo or research group allocation email us and cc the PI who will authorize it.

Different Node Types

- Login nodes
    - Four virtual machines
    - This is where you are when you log in
    - No heavy computation, interactive jobs, or long running processes
    - Script or code editing
    - Job submission
- Compile nodes
    - Where you compile code
    - Explore the Summit software environment
- Compute/batch nodes
    - This is where jobs that are submitted through the scheduler run
    - Intended for heavy computation

Storage Spaces

- **Home Directories**
  - /home/$USER
  - Not for direct computation
  - Small quota (2 GB)
  - Backed up

- **$PROJECT Space**
  - /projects/$USER
  - Mid level quota (250 GB)
  - Large file storage
  - Backed up

- **Scratch Directory**
  - /scratch/summit/$USER
  - 10 TB
    - Can ask for more if needed
  - Files purged around 90 days

What Belongs Where?

- /home
  - Scripts
  - Code
  - Small, important files/directories
  - Inappropriate for sharing files with others
  - Inappropriate for job output

- /projects
  - Code/files/libraries
  - Software you are installing
  - Mid-level size input files
  - Appropriate for sharing files with others
  - Inappropriate for job output

- /scratch/summit
  - Output from running jobs
  - Large files/datasets
  - Appropriate for sharing files with others
  - THIS IS NOT APPROPRIATE FOR LONG TERM STORAGE

Transferring Data

- Globus is Research Computing's preferred method of data transfer

- Designed with researchers in mind

- End points between computers make for efficient data transfer with an easy to use interface
  - Endpoints are different locations that data can be moved to/from
  - Personal or multi-user

- rsync and sftp through the login nodes is good for small transfers – transfers that take a few minutes.

Setting Up Globus

- Create an account at Globus.org

- Make your personal computer an endpoint if you want
  - Not needed if you are transferring between two other endpoints, like a repository and RC

- Transfer data
  - www.globus.org

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Software

- Common software is available to everyone on the systems
- Can install your own software
  - But you are responsible for support
  - We are happy to assist
- Research Computing uses modules to manage software
  - You load modules to prepare your environment for using software
    - Modules set any environment variables, paths,etc.
    - Set environment so application can find appropriate libraries, etc.

Important Things to Know About Modules

- Some modules might require a specific hierarchy to load
    - For some modules, you may need to specify a specific version
        - For example, `module load R/3.3.0`
    - For other modules, you may be able to be more generic
        - For example, `module load matlab`

- Some modules may require you to first load other modules that they depend on
- To find dependencies for a module, type `module spider <package>`

- To find out what software is available, you can type **`module avail`**

- To set up your environment to use a software package, type **`module load <package>/<version>`**

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Job Submission

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Running Jobs

- What is a "job"?
  - Jobs are scripted packages of work for the cluster to perform
  - Have a job number

- Interactive jobs
  - Work interactively at the command line of a compute node

- Batch jobs
  - Submit job script which will be executed when resources are available
    - Create a text file, shell script containing information about the job
    - Submit the job file to a queue

Job Scheduling

- On a supercomputer, jobs are scheduled rather than just run instantly at the command line
  - Shared system
  - Jobs are put in a queue until resources are available

- Need software that will distribute the jobs appropriately and manage the resources
  - Simple Linux Utility for Resource Management (Slurm)
    - Keeps track of what nodes are busy/available, and what jobs are queued or running
    - Tells the resource manager when to run which job on the available resources

Research Computing
UNIVERSITY OF COLORADO BOULDER

Be Boulder.

Partitions and 'Quality of Services'

- Defining where and how your job will run

- Partitions (basically a queue):
  - Resources/hardware

- QoS:
  - Tells what the limits or characteristics of a job should be
    - Maximum wall time (timelimit)
    - Number of nodes

- One partition might have multiple QoS

- A QoS might exist on multiple partitions

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Available Partitions

Partition	Description	# of nodes	cores/node	GPUs/node
shas	General Compute (Haswell)	380	24	0
sgpu	GPU-enabled nodes	10	24	effectively 4
smem	High-memory nodes	5	48	0
sknl	Phi (Knights Landing) nodes	20	68	0

Quality of Service

QoS	Description	Maxwall	Max jobs/user	Max nodes/user
normal	Default QoS	Derived from partition	n/a	256
testing	For quick turnaround when testing	30 Min	1	2/user; max 12 cores/node
interactive	For interactive jobs (command or GUI)	4 Hours	1	1 core
long	For jobs needing longer wall times	7 D	n/a	20
condo	For groups who have contributed to the Summit condo	7 D	n/a	n/a

Useful Slurm Commands - sbatch

- **sbatch**: submit a batch script to slurm
- You can use a bunch of flag options in a batch script or on the command line
- Useful to put in script so have for future use

- Example:

```
sbatch test.sh
```
OR

```
sbatch --partition=shas test.py
```

http://slurm.schedmd.com/sbatch.html

SBATCH Options

http://slurm.schedmd.com/sbatch.html

```
#SBATCH <options>        sbatch <options>
```

- Allocation: `--account=<account_no>`
- Partition: `--partition=<partition_name>`
- Sending emails: `--mail-type=<type>`
- Email address: `--mail-user=<user>`
- Number of nodes: `--nodes=<nodes>`
- Number of tasks: `--ntasks=<processes>`
- Quality of service: `--qos=<qos>`
- Reservation: `--reservation=<name>`
- Wall time: `--time=<wall time>`
- Job Name: `--job-name=<jobname>`

- FYI: You do NOT actually type <> above – this designates something specific you as a user must enter about your job

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

New User Seminar   38

Be Boulder.

Working on Summit

- If you use both Summit and Blanca, make sure you load the appropriate slurm module – if on a Login Node:

```
$ module load slurm/blanca
$ module load slurm/summit
$ ml slurm/summit #shorthand
```

- If on an scompile node, this is not needed and will return an error.

- After you run this command you can run sbatch to submit jobs

Practice Examples

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Submit Your First Job!

- Submit a slurm job with the following instructions:

1. The job should run the Unix "hostname" command
2. The job will be submitted from a bash script named hostname_summit.sh
3. The job will run on 1 node
4. We will request 1 minute wall time
5. Run using "testing" QOS
6. Run on the shas-testing partition

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Hostname_summit.sh

```bash
#!/bin/bash
#SBATCH --nodes=1                       # Number of requested
nodes
#SBATCH --time=0:01:00                  # Max wall time
#SBATCH --qos=testing                   # Specify debug QOS
#SBATCH --partition=shas-testing        # Specify Summit
haswell nodes
#SBATCH --output=hostname_%j.out        # Rename standard
output file


# purge all existing modules
module purge


hostname
```

Research Computing
UNIVERSITY OF COLORADO BOULDER

Be Boulder.

Running the script

- Load up the slurm module (probably not needed)
  ```
  module load slurm/summit
  ```

- Submit the job:
  ```
  sbatch hostname_summit.sh
  ```

- Check output

Another slurm command

- **squeue**
  - View information about jobs located in the slurm scheduling queue
- OPTIONS:
  - User: `-u <user_list>`
  - Queues: `--qos=<qos_list>`

- EXAMPLE:

  `squeue --qos=debug`

http://slurm.schedmd.com/squeue.html

Your turn

- Submit a slurm job with the following instructions:

1. The job should run first the whoami command, then the Unix "sleep" command for 30 seconds, then the hostname command
   - Syntax for these Unix commands are below:

```
whoami
sleep 30
hostname
```

http://slurm.schedmd.com/squeue.html

Your turn

• Submit a slurm job with the following instructions:

1. The job will be submitted from a bash script named sleep.sh
2. The job will run on 1 node
3. Request a 1 minute wall time
4. Run the job from the normal QOS
5. Run the job from the Summit haswell partition
6. Name your job sleep
7. Email yourself the results at the end of the job run
   • Hint: This requires two SBATCH directives to do this – see link at top of this slide

Sleep.sh

```bash
#!/bin/bash
#SBATCH --nodes=1                              # Number of requested nodes
#SBATCH --time=0:01:00                         # Max walltime
#SBATCH --qos=normal                           # Specify normal QOS
#SBATCH --partition=shas                       # Specify Summit GPU nodes
#SBATCH --output=sleep_%j.out                  # Rename standard output file
#SBATCH --job-name=sleep                       # Job submission name
#SBATCH --mail-type=end                        # Email you when the job ends
###SBATCH --mail-user=<user>@colorado.edu      # Email address to send to


# purge all existing modules
module purge


whoami
sleep 30
hostname
```

Running an external script

- Let's run a Matlab program

- We will run the bash script matlab.sh

- This script calls and runs matlab_tic.m

Running the script

- Submit the job:
    ```
    sbatch matlab.sh
    ```


- Check output

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Matlab.sh

```bash
#!/bin/bash
#SBATCH --nodes=1                          # Number of requested nodes
#SBATCH --time=0:02:00                     # Max walltime
#SBATCH --qos=testing                      # Specify debug QOS
# #SBATCH --partition=shas                 # (disabled) specify Summit haswell nodes
#SBATCH --output=matlab_%j.out             # Output file name

# purge all existing modules
module purge
# Load a specific Matlab module
module load matlab/R2016b
# Run matlab without a GUI
matlab -nodisplay -nodesktop -r "clear; matlab_tic;"
```

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Your turn

- Submit a slurm job with the following instructions:

1.   Create an R program called `R_program.R` that creates a vector called "planets" and then list the planets in the vector
   - Syntax: `planets -> planets <- c("Mercury", "Venus", "Earth", "Mars", "Jupiter", "Saturn", "Uranus", "Neptune", "Pluto")`
2.   Print off the vector
   - Syntax: `planets`
3.   Create a bash script called `R_code.sh` that runs the R script
   - Syntax: `Rscript R_program.R`
4.   The job will run on 1 node
5.   We will request a 1 minute wall time
6.   Specify the debug QOS
7.   Specify the shas partition
8.   The output will be put in a file called R_code_%j.out
9.   Don't forget to load the R module!

Solution – R_code.sh

```bash
#!/bin/bash
#SBATCH --nodes=1                      # Number of requested nodes
#SBATCH --time=0:01:00                 # Max walltime
#SBATCH --qos=debug                    # Specify debug QOS
#SBATCH --partition=shas               # Specify Summit haswell nodes
#SBATCH --output=R_code_%j.out         # Output file name


# purge all existing modules
module purge


# Load the R module
module load R/3.3.0


# Run R Script
Rscript R_program.R
```

Research Computing
UNIVERSITY OF COLORADO BOULDER

Be Boulder.

Solution – R_program.R

```r
#Simple R code example by Shelley Knuth
(shelley.knuth@colorado.edu)


# Create vector

planets <- c("Mercury", "Venus", "Earth", "Mars",
"Jupiter", "Saturn", "Uranus", "Neptune", "Pluto")


# Print off vector

planets
```

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Interactive jobs!

- Sometimes we want our job to run in the background
- Sometimes we want to work in program in real time
- For example, Matlab
- Let's run an interactive Matlab job

Interactive job using X

- In this example we'll use X windows to access the Matlab GUI
- To do this, we are going to log out and log back in
  - Only necessary for demo
  - Need to add something to the sign in process

- For Mac Users:
  ```
  ssh -X username@login.rc.colorado.edu
  ```

  - For Windows Users, must set up X-forwarding through your SSH client program
- Also must have an X-server package on your laptop
  - Xming for Windows or XQuartz for Mac

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Interactive job

- To work with Matlab interactively, we're going to request some time from the supercomputer
- When the resources become available then we will start up Matlab
- Commands to run:
  ```
  module load slurm/summit

  sinteractive --qos=debug --time=00:05:00
  ```

  Once we receive a prompt, then:
  ```
  module load matlab
  matlab
  ```

- Once we finish we must exit!

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.

Questions?

- Email rc-help@colorado.edu
- Twitter: @CUBoulderRC

- Link to survey on this topic:
  http://tinyurl.com/curc-survey18

- Slides:
  https://github.com/ResearchComputing/New_User_Seminar

Research Computing
UNIVERSITY OF COLORADO **BOULDER**

Be Boulder.