



# New User Seminar



Research Computing  
UNIVERSITY OF COLORADO **BOULDER**

**Be Boulder.**

# Research Computing New User Seminar

Instructor:

- Website: [www.rc.colorado.edu](http://www.rc.colorado.edu)
- Helpdesk: [rc-help@colorado.edu](mailto:rc-help@colorado.edu)
- Slides: [https://github.com/ResearchComputing/New\\_User\\_Seminar](https://github.com/ResearchComputing/New_User_Seminar)
- Survey: <http://tinyurl.com/curc-survey18>

# Account Check

Does anyone ***not*** have a **CU Research Computing account** who would like to use a temporary account\*?

*\*only available during seminar*

# Learning Goals

1. What is CU Research Computing (CURC)?
2. Understand Basic Resources
3. Getting an account & logging in
4. Navigate the RC system
5. Run a job
6. Help!

# Things to Take Note Of

- Confusing, ambiguous, highly nuanced concepts
- Goal is to avoid common mistakes, pitfalls, and frustrations



**Ask Questions!**

# CU Research Computing provides

Large-scale, beyond-the-desktop  
computing resources for the  
University of Colorado and more!




Colorado  
State  
University




# Resources include:

- High Performance Computing
- Storage of Research Data
- High-Speed Data Transfer
- Data Sharing
- Cloud Computing
- Training & Education





# Primarily known for: High Performance Computing (HPC)






# High Performance Computing (HPC) vs. Traditional Computing

- Traditional computing *generally* has access to a single processor (perhaps multiple cores)




# High Performance Computing (HPC) vs. Traditional Computing

- HPC is a network (or cluster) of hardware we call “nodes” linked together with high-speed interconnects



# High Performance Computing (HPC)

- HPC is a multiprocessor environment that allows users to run jobs on several processors at once (also called parallel processing)



# What can / use HPC for?

- Solving large problems that require more:
  - Memory than you have on your PC
  - cores/nodes/power than you have on your PC
- Jobs that require hardware you may not have:
  - High Performance GPU computing
  - Specific Operating System
- Visualization rendering

# Puzzle Analogy

- Working on a puzzle by yourself (although enjoyable!) will take time, even if you're fast!



- You can speed it up by inviting friends over and enticing them with food




14

- What happens if there are too many people?




# What can / use HPC for?

- Jobs that would take a long time on local machines can instead be distributed over hardware:
  - Parallelized to split up then joined (if software enabled)




- Broken up into many serial jobs




# Research Computing Resources

# HPC Clusters at CURC


**Summit**



**Blanca**




## Summit



- NSF-Funded
- Shared
- 450+ Nodes (mostly Intel Xeon Haswell)
  - 24 cores per “shas” (general compute) node
  - Additional Node types:
    - 10 GPU Nodes (NVIDIA Tesla K80)
    - 5 High Memory Nodes (2TB ram)
    - 20 Intel Xeon Phi nodes
- **11,400 total cores**

## Blanca



- Buy-in Cluster
  - Per user specs and vendor constraints
- High priority use on your hardware
  - Preemptable use to all other nodes
- Heterogenous
  - ~223 Compute Nodes
  - ~10 GPU nodes


# Storage at CURC

## Core




- Included with RC account
  - /home
  - /projects
  - scratch space

## PetaLibrary



- Paid Service for:
  - Storage
  - Archive
  - Sharing of research data

## Local or Cloud




- You can download your data locally or to a variety of other cloud resources




# Data Sharing: Within RC

- Sharing workspaces
  - Project space
  - Scratch space
  - PetaLibrary space



# Data Sharing: Outside RC

- Globus (recommended):
  - GUI Web Application
  - Automates large transfers
  - Resumes failed transfers
  - Distributes large transfers across DTNs
  - Endpoints that can shared
- Data Transfer Nodes (DTN)
  - Internal CU network needed
- Command line tools
  - scp, sftp, rsync, rclone



# Accessing Research Computing

# How to Access RC Resources?

1. Get an account
2. Set up two factor authentication
3. (Inform us of any specific needs)
4. Log in
5. Create greatness! (responsibly)

# Getting an account

- CU Boulder users and affiliates:
  - Request an account through the RC Account request portal
  - <https://rcamp.rc.colorado.edu/accounts/account-request/create/organization>
- CSU Users:
  - Request an CSU eID if you don't have one
  - Fill out account application form
  - Duo authentication
  - Then get an RC user account
  - <https://www.acns.colostate.edu/hpc/summit-get-started/>
- RMACC Users:
  - Contact your local representative, if known. Email [rc-help@colorado.edu](mailto:rc-help@colorado.edu)
  - We'll guide you through the process

# Demo: Getting an Account

- CU Boulder users and affiliates:
  - Request an account through the RC Account request portal
  - <https://rcamp.rc.colorado.edu/accounts/account-request/create/organization>



# Your RC Account

## Access to:


1. Summit Cluster (soon Alpine)
2. “Core” storage
3. PetaLibrary Storage (if applicable)
4. Gateways (CU affiliates)

# Two Factor Authentication

- Provides an extra level of authentication
  - We are outside the firewall!
  - Valuable resources
  - Inviting, high-profile target
  - Lost time investigating/fixing
- Duo
  - You will receive a Duo invitation when your RC account is created.

# Duo Authentication

1. Most users use the Duo smartphone app
2. “Phone Call” is an alternatives
3. Physical code generator “token” available for \$20



# Linux comfort level check!

- On a scale from (1-10) **how familiar/comfortable are you with Linux?**
  - The command line
  - Basic commands
  - Linux filesystem
  - Navigating the filesystem

# Terminal

- Mac or Linux
  - Terminal application
- Windows
  - PuTTY (recommended)
  - Powershell

```
[user0083@tlogin1 ~]$ pwd  
/home/user0083  
[user0083@tlogin1 ~]$ █
```

- Open OnDemand (*alternative for CU affiliates*)
  - For those less familiar with Linux ([ondemand.rc.colorado.edu/](https://ondemand.rc.colorado.edu/))

# Demo: Logging in via Terminal

- To login to an RC login node:

```
ssh <username>@login.rc.colorado.edu
```

Supply your IdentiKey password and your Duo app will alert you to confirm the login

---

If you're using a tutorial account:

```
ssh <username>@tlogin1.rc.colorado.edu
```



# Demo: Logging in with OnDemand

---

# Logging In

- It's important to note that you are NOT logging into any specific resource, Summit, etc.
- When you log in, you land on our login nodes
- From there, you can access our other resources



# Navigating Research Computing

## Node

- One computing server
- Physical hardware
- Work together in parallel

## File System

- The basic tree-like layout
- From most nodes\* you have access to the entire file system





## Node

- One computing server
- Physical hardware
- Work together in parallel

## File System

- The basic tree-like layout
- From most nodes\* you have access to the entire file system



# Node Types

Login	Compile	Compute
<ul style="list-style-type: none"><li>• Where you start</li><li>• For editing code, job submission</li><li>• No heavy computation</li></ul>	<ul style="list-style-type: none"><li>• Where you compile code, install packages</li><li>• Explore the Summit software environment</li><li>• Edit code, submit jobs</li><li>• No heavy computation</li></ul>	<ul style="list-style-type: none"><li>• Where scheduled jobs run</li><li>• Intended for heavy computation</li></ul>
Ex. edit job script	Ex. Install python libs	Ex. Running Matlab

# Demo: Exploring nodes

- Once logged in:

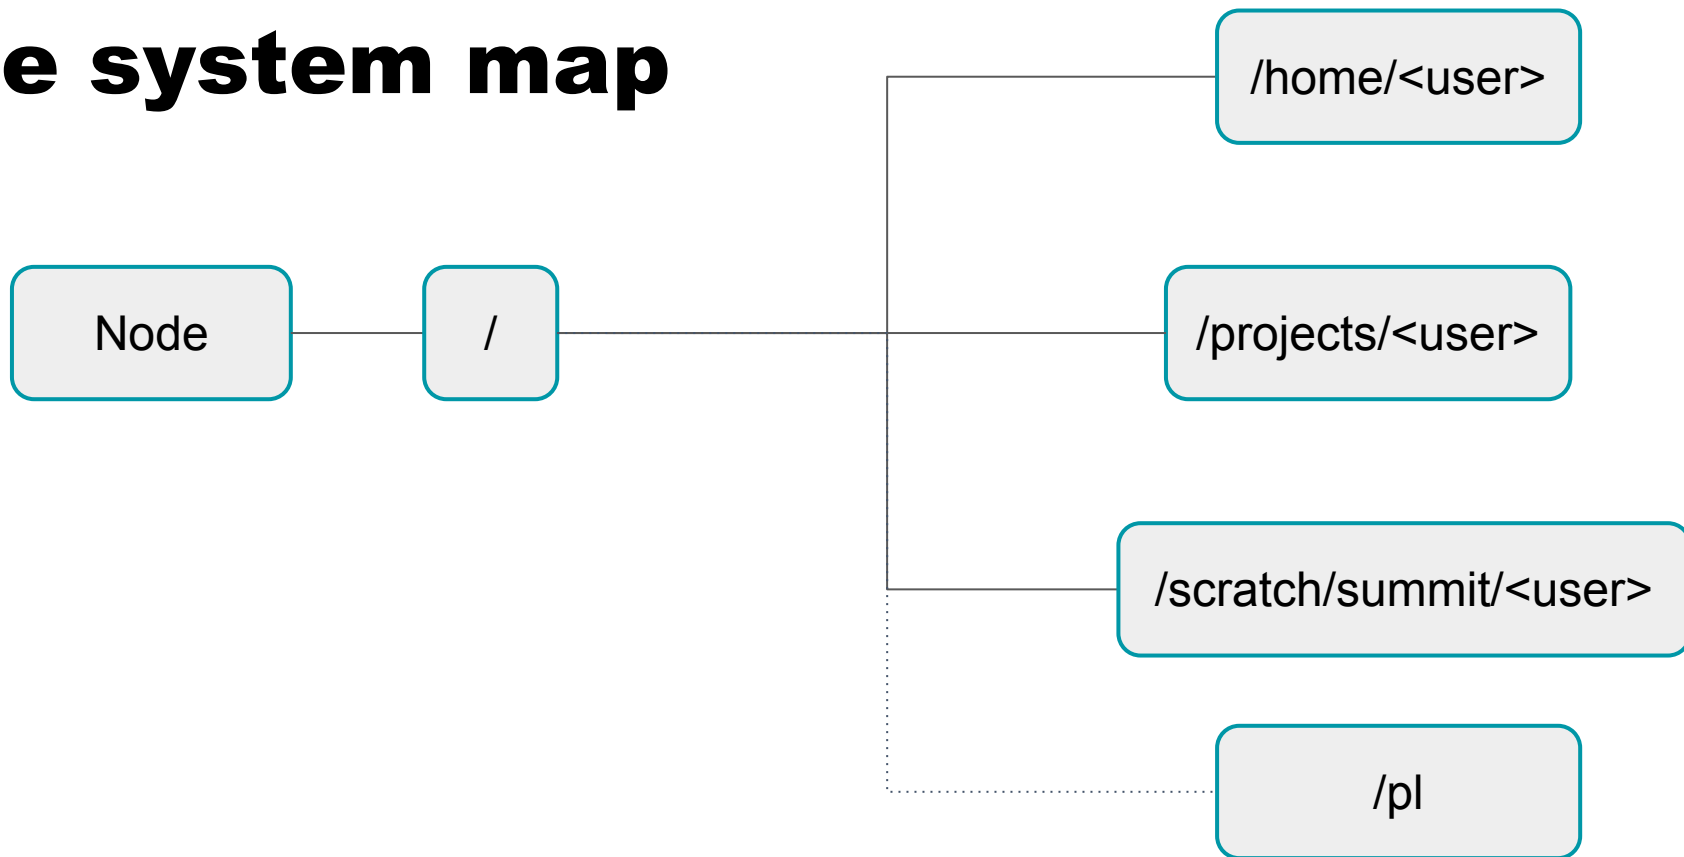
```
$ ssh scompile
```

To log in to a compile (or head) node.

```
$ module avail
```

To check currently available software

# File system map





# File Structures

/home (2GB)	/projects (250GB)	Scratch (10TB)
<ul style="list-style-type: none"><li>• Scripts, Code, Small, important files/directories</li><li>• Not for sharing files or job output</li></ul>	<ul style="list-style-type: none"><li>• Code/files/libraries</li><li>• Software you are installing</li><li>• Sharing files</li><li>• Not for job output</li></ul>	<ul style="list-style-type: none"><li>• Output from running jobs</li><li>• Large files/datasets</li><li>• Sharing files</li><li>• Not for long term storage</li></ul>
Ex .bashrc	Ex. Shared job scripts	Ex. Data

# Demo: Exploring the Filesystem

- Once logged in:

```
cd /home/<user>
```

```
cd /projects/<user>
```

```
cd /scratch/summit/<user>
```

To navigate to your different directories

# Using Research Computing

- We have logged on
- We have explored nodes and filesystem
- But how do we actually *use* the computing resources?

# Running a Job

# The fundamental “job”

What is a “*job*”?

- Work for the cluster to perform on

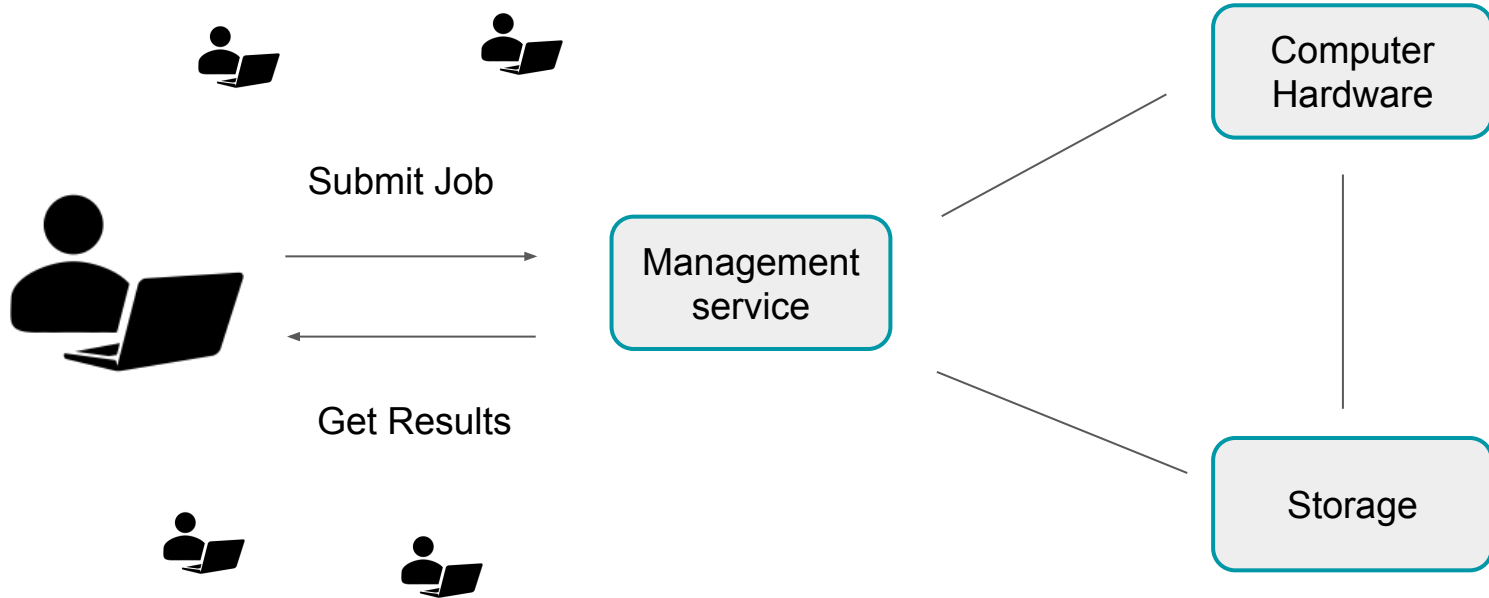
## 1. Batch jobs

- Submit job script which will be executed when resources are available
  - Create script containing information about the job
  - Submit the job file to a queue


## 2. Interactive jobs

- Work interactively at the command line of a compute node

# Job Scheduling




# Job Scheduling



# Your first job

- Where to write it?
- How to write it?
- How to run it?





# Job Script: 3 main parts

## 1. Directives

- Specify resource requirements

## 2. Software

- Because jobs run on a different node than from where you submitted...
- ...software that is needed must be loaded via the job script

## 3. User scripting

- the actual user scripting that will execute when the job runs

```
#!/bin/bash
```

```
## Directives
```

```
#SBATCH --<resource>=<amount>
```

```
## Software
```

```
module purge
```

```
## User Scripting
```

# Directive Options

<http://slurm.schedmd.com/sbatch.html>

```
#SBATCH <options>      sbatch <options>
```

- Allocation: `--account=<account_no>`
  - Partition: `--partition=<partition_name>`
  - Sending emails: `--mail-type=<type>`
  - Output file: `--output=<file name>` (%j gives you job id)
  - Number of nodes: `--nodes=<nodes>`
  - Number of tasks: `--ntasks=<processes>`
  - Quality of service: `--qos=<qos>`
  - Reservation: `--reservation=<name>`
  - Wall time: `--time=<wall time>`
  - Job Name: `--job-name=<jobname>` ...etc...
- 
- FYI: You do NOT actually type <> above – this designates something specific you as a user must enter about your job

# Writing a Job: `hostname.sh`

- Submit a slurm job with the following instructions:
  1. The job will be submitted from a bash script **named `hostname_summit.sh`**
  2. The job will run **on 1 node**
  3. Request **1 minute wall time**
  4. Run on the **`shas-testing` partition**
  5. Output should **contain job id**
  6. The job should run **the Unix “`hostname`” command**

# Writing a Job: Hostname

- Set up job boilerplate (nano or vim):

```
#!/bin/bash

## Directives
#SBATCH --<resource>=<amount>

## Software
module purge

## User Scripting
```

# Writing a Job: Hostname

- Directives:
  - 1 node
  - 1 minute wall time
  - “shas-testing” partition
  - Output should contain job id
- Number of nodes:  
`--nodes=<nodes>`
- Wall time:  
`--time=<wall time>`
- Partition:  
`--partition=<partition_name>`
- Output file:  
`--output=<file name>`  
(%j gives you job id)

# Writing a Job: Hostname

- Software:
  - Do we need any software? (modules)
- User script:
  - What command do we want to run?

# Demo: Writing Hostname

# Demo: Writing Hostname

```
#!/bin/bash
#SBATCH --nodes=1                # Number of requested nodes
#SBATCH --time=0:01:00          # Max wall time
#SBATCH --partition=shas-testing # Specify Summit haswell nodes
#SBATCH --output=hostname_%j.out # Rename standard output file

# purge all existing modules
module purge

hostname
```



# Submitting a Job

1. Load up the slurm module (default)

```
module load slurm/summit
```

2. Submit the job:

```
sbatch <script name>.sh
```

3. Check output

# **Review: Learning Goals**

- 1. Why High Performance Computing (HPC)?**
- 2. Understand Basic Resources**
- 3. Get an account & log in**
- 4. Navigate the RC system**
- 5. Run a job**
- 6. Help!**

# Help! I'm stuck, where do I go?

- **Documentation**: [curc.readthedocs.io/](https://curc.readthedocs.io/)
- **Trainings with Center for Research Data and Digital Scholarship (CRDDS)**: <https://www.colorado.edu/crdds/>
- **Helpdesk**: [rc-help@colorado.edu](mailto:rc-help@colorado.edu)

# Questions

# Survey and feedback

<http://tinyurl.com/curc-survey18>